

# **Parameter-Efficient Fine-tuning**

# Parameter-Efficient Fine-tuning

- A method of pre-training models by changing a small number of parameters
- Used for training LLMs that are too costly to fine-tune
- When training on several different tasks, you can store only the set of changed parameters for each task

# Few-shot and Zero-shot for GPT

- GPT-3 has been trained on 45 TB of data. This gives it useful properties.
- The model can be applied to a new problem by showing it examples of solutions in a prompt.

Translate from Russian to English:

стол => table

сыр => cheese

дом =>

It's called **Few-shot** learning.

# Few-shot and Zero-shot for GPT

It is possible not to show the correct solutions at all

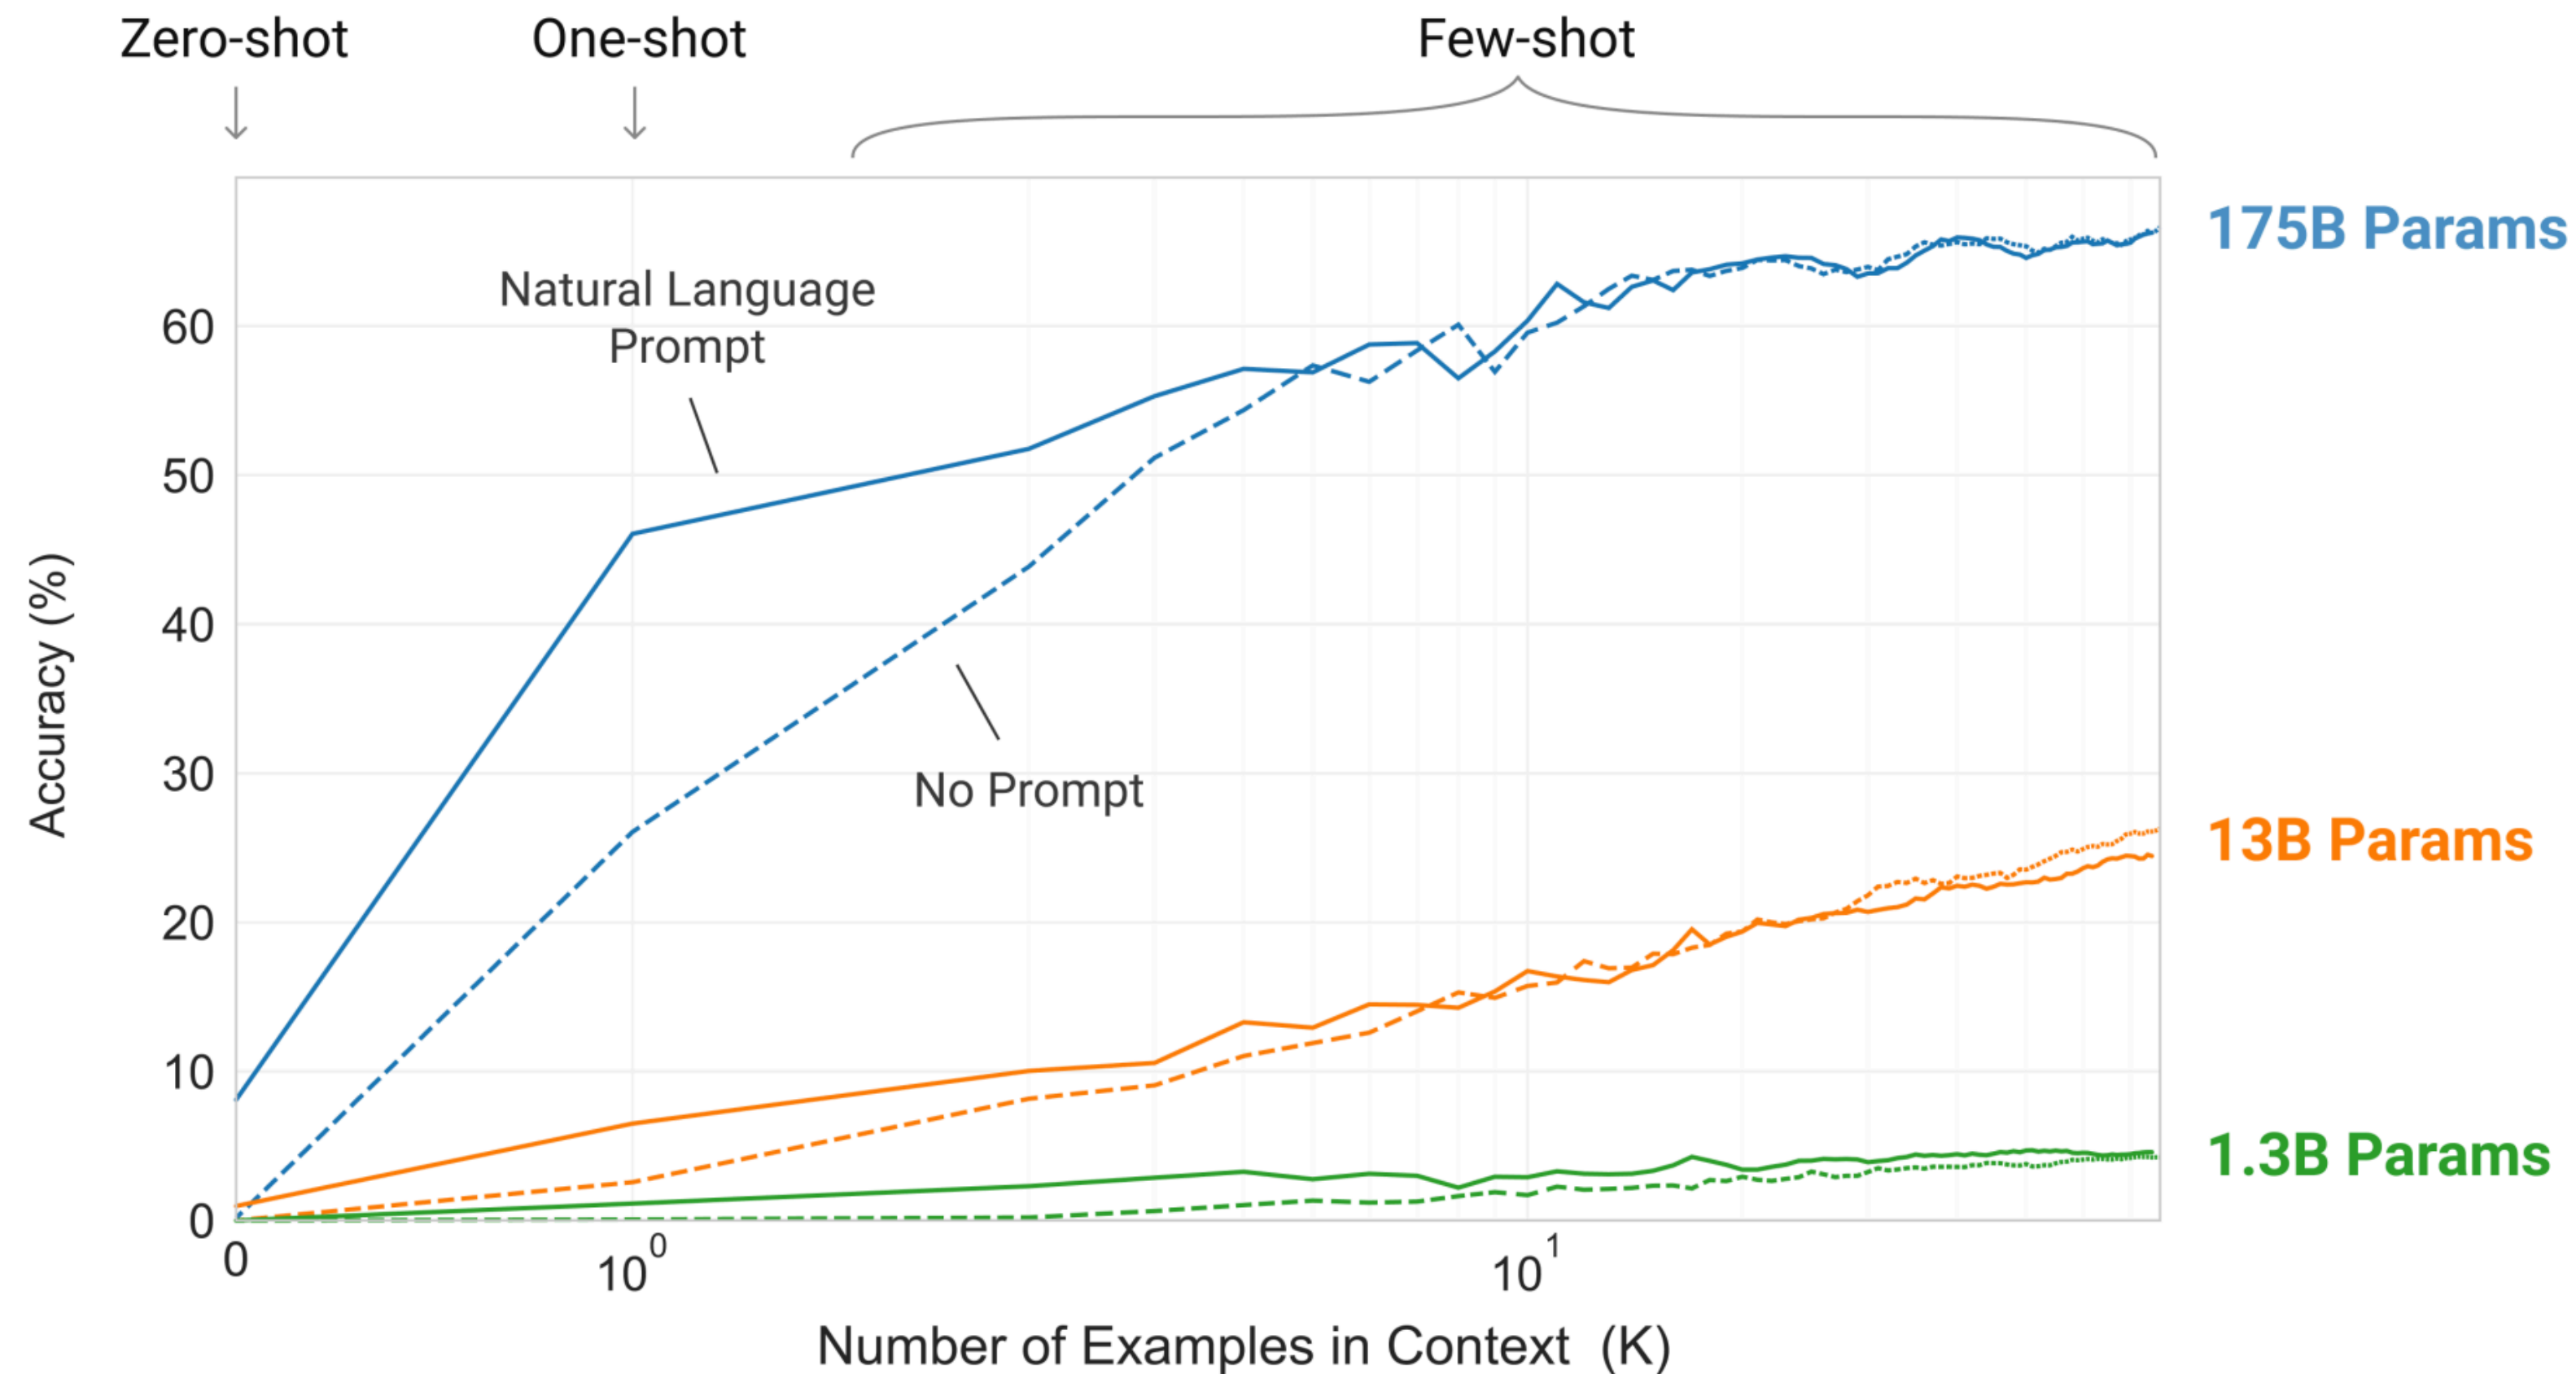
Translate from Russian to English:

ДОМ =>

It's called **Zero-shot** learning.

# Few-shot and Zero-shot for GPT

The larger the model, the better it performs.



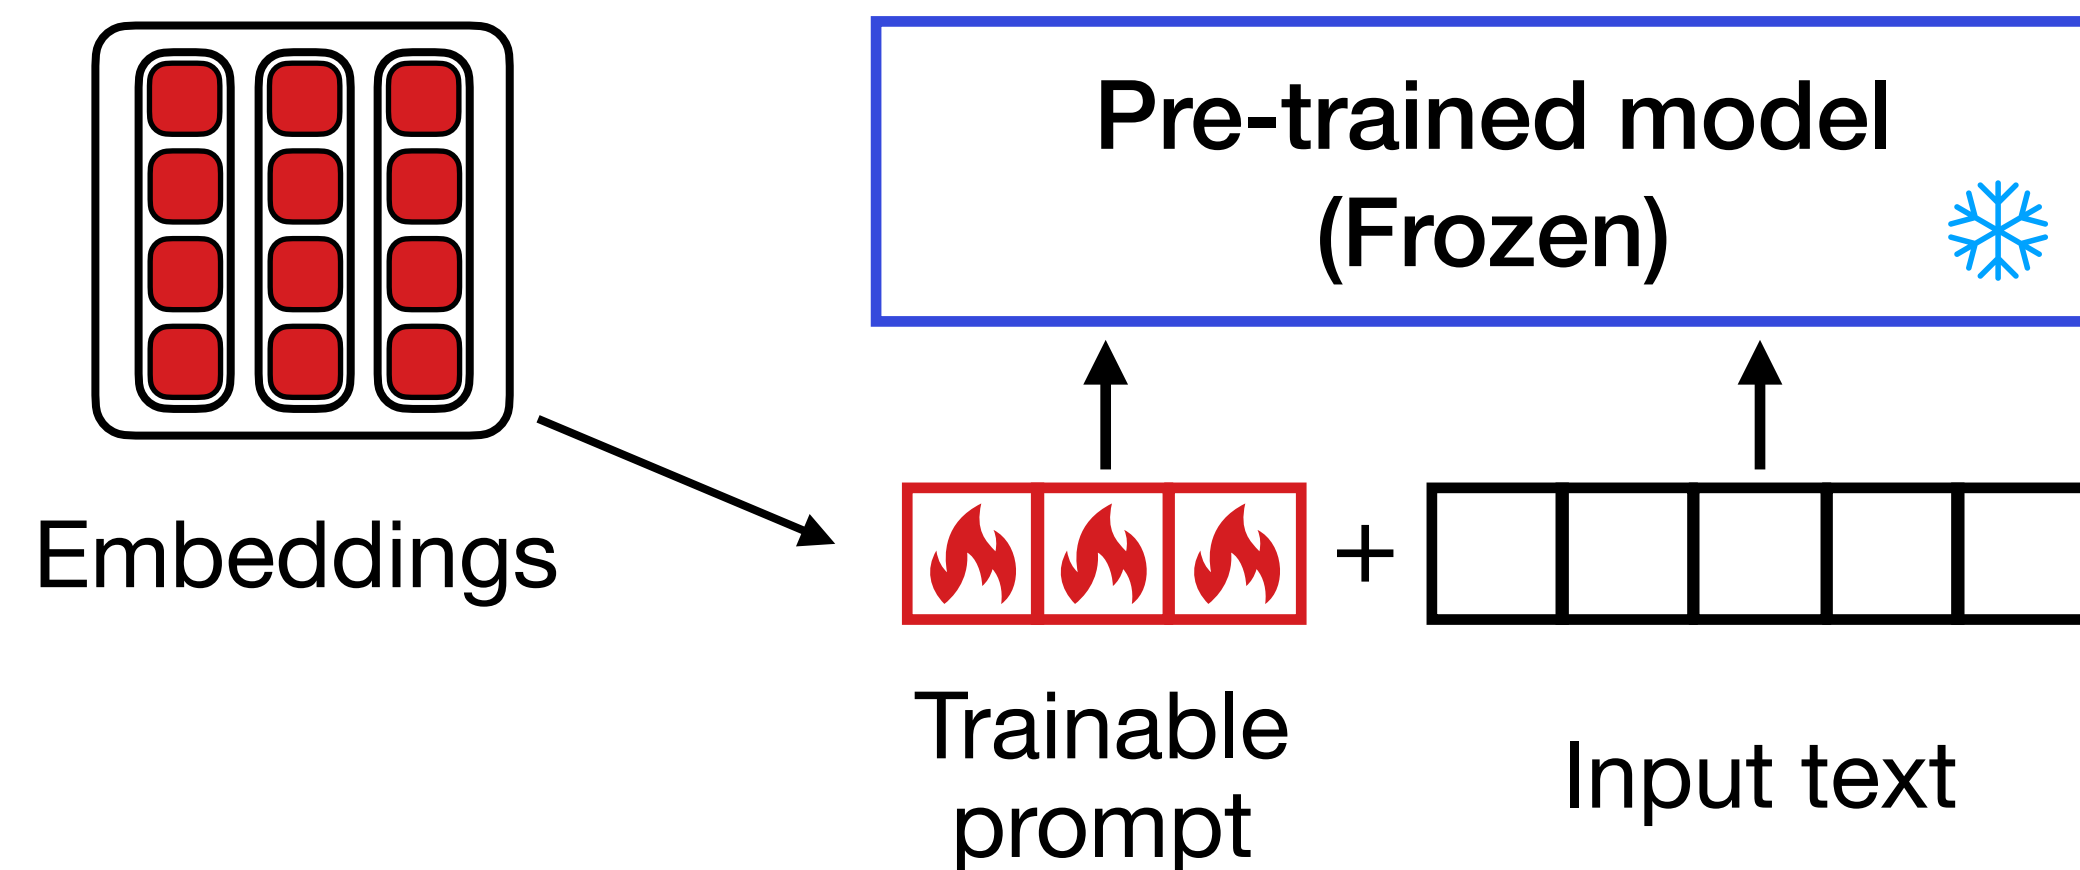
However, the result may be poor because of the fact that

- The model has not learnt how to solve this problem
- Prompt is not good enough

# Prompt Tuning

**Idea:** Let's try to automatically select the most appropriate prompt

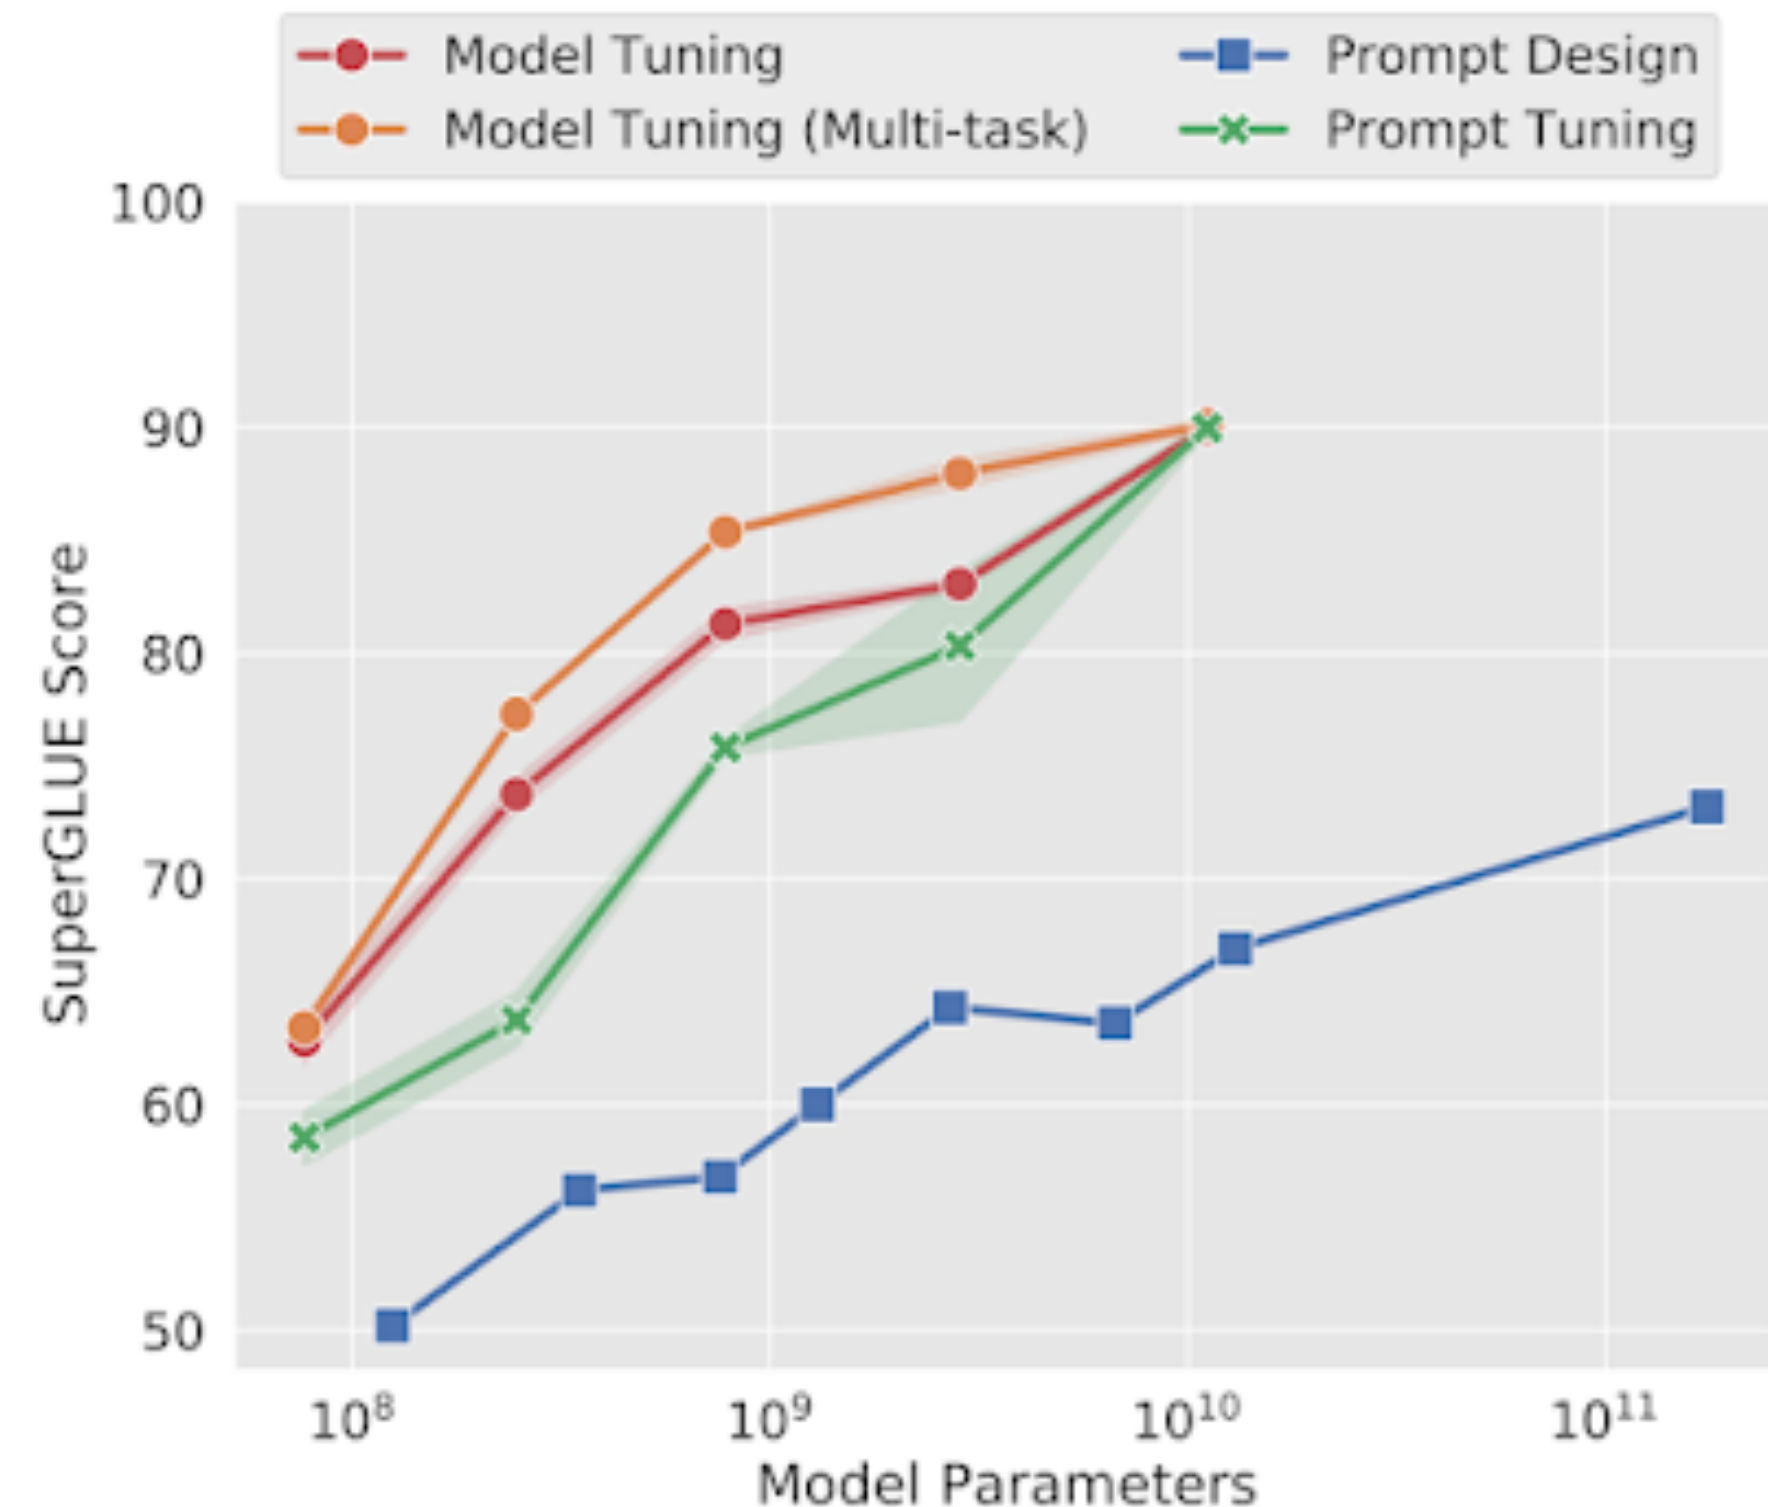
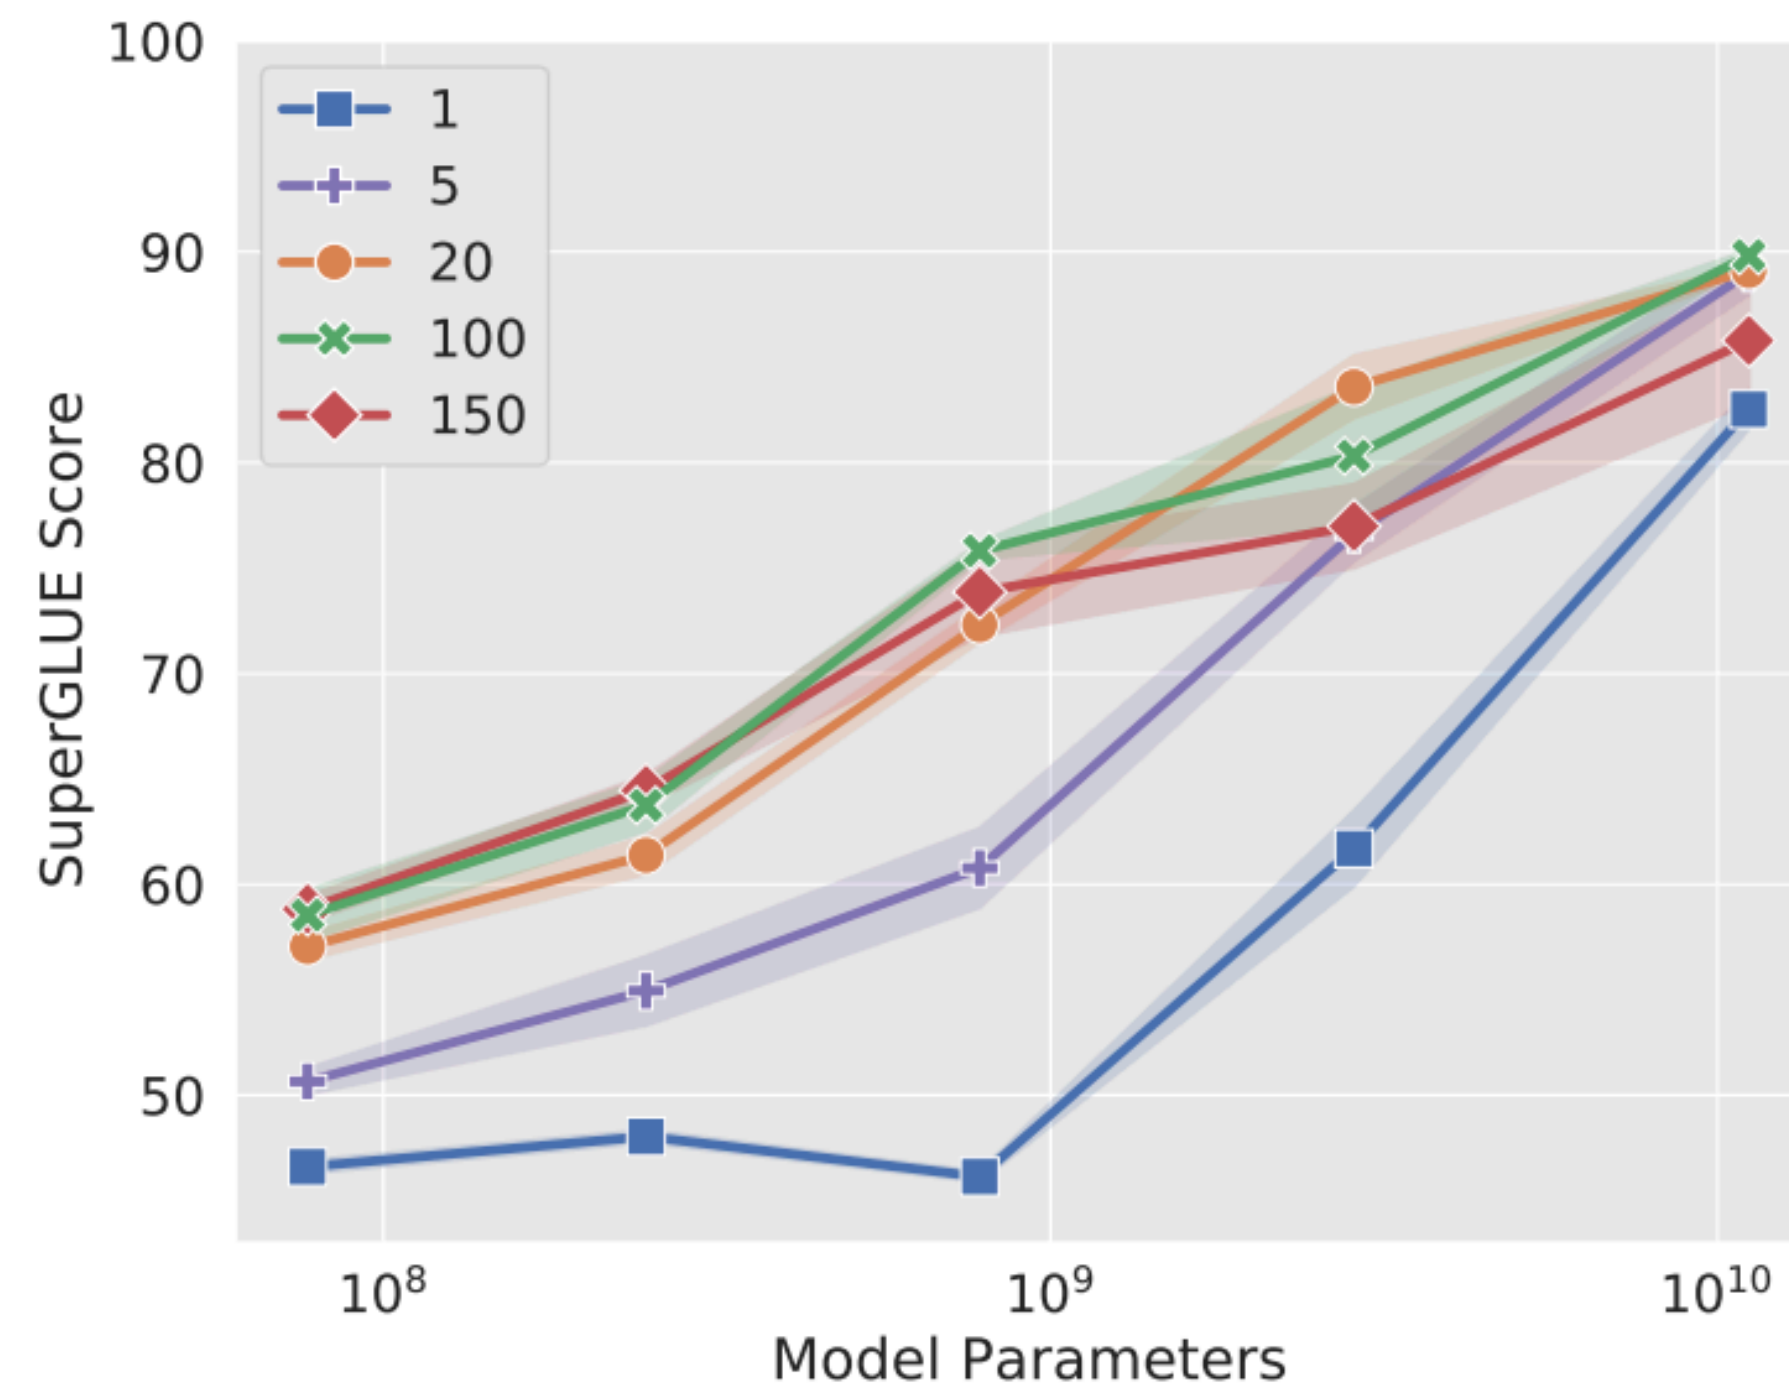
- Initialise the prompt embeddings randomly, specifying only the number of them
- We can initialize the embeddings with embeddings of some prompt.
- For the classification task we need to train the head additionally



# Prompt Tuning

- The length of the prompt directly affects the quality of the model
- However, the larger the model, the smaller the difference in quality

Prompt length



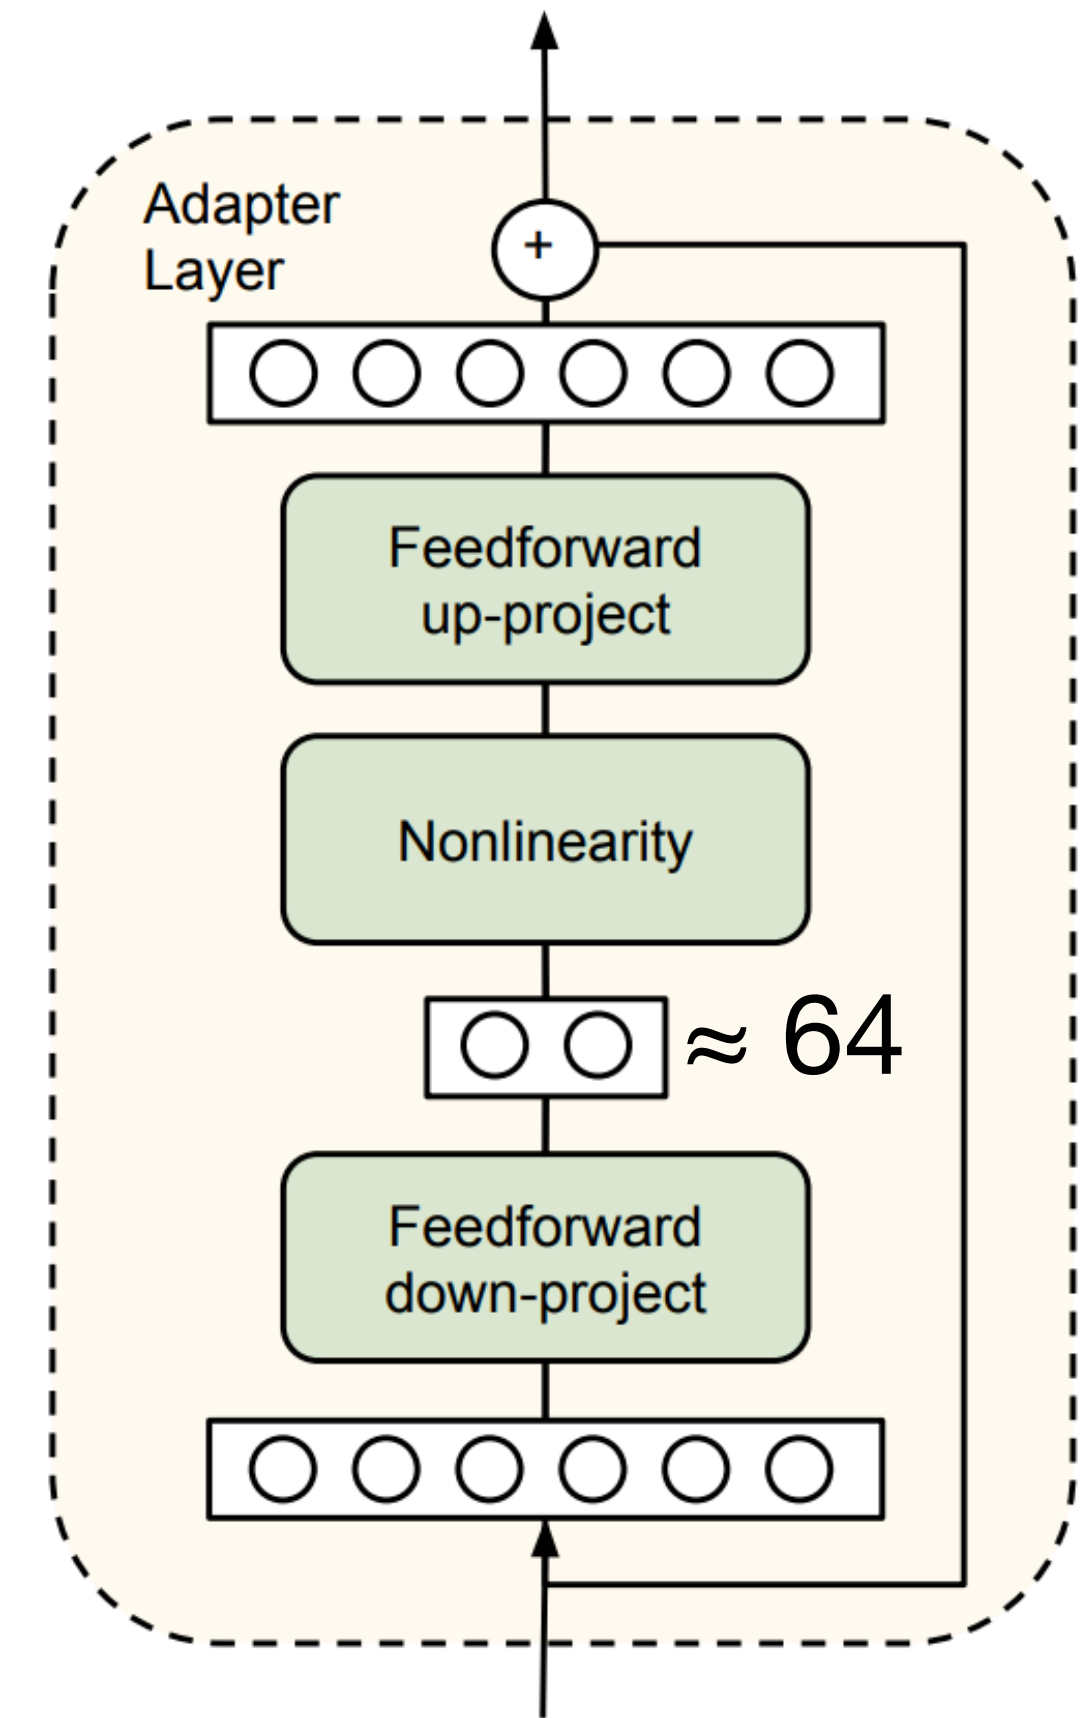
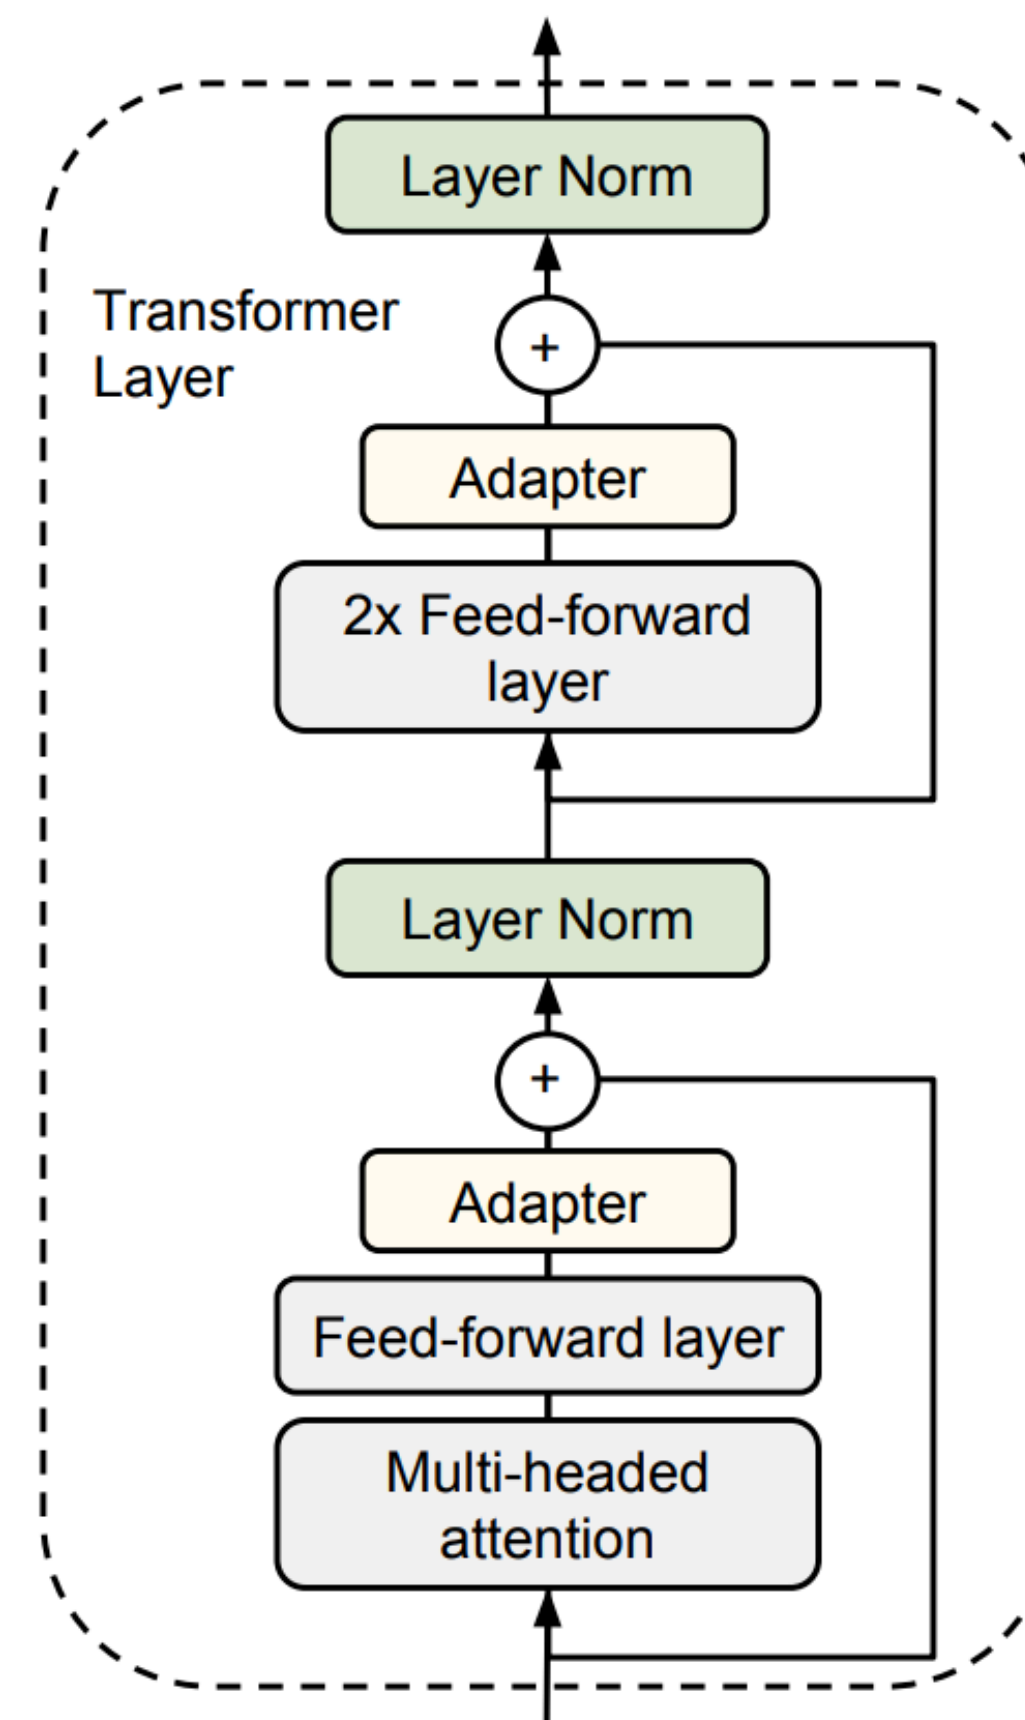
# Prompt Tuning: Disadvantages

- Addition of a trained prompt limits the maximum length and slows down the model
- Prompt Tuning is very unstable and quality changes non-monotonically as the size of the prompt and model increases



# Adapters

- After each attention layer and FFN, a **small** trainable adapter is added
- The adapter has skip-connection and two full-connection layers with dimensionality reduction (reduces the number of parameters)
- Only the adapters, normalizations and head are trained
- Such technique reaches fine-tuning in terms of quality



# Adapters: Disadvantages

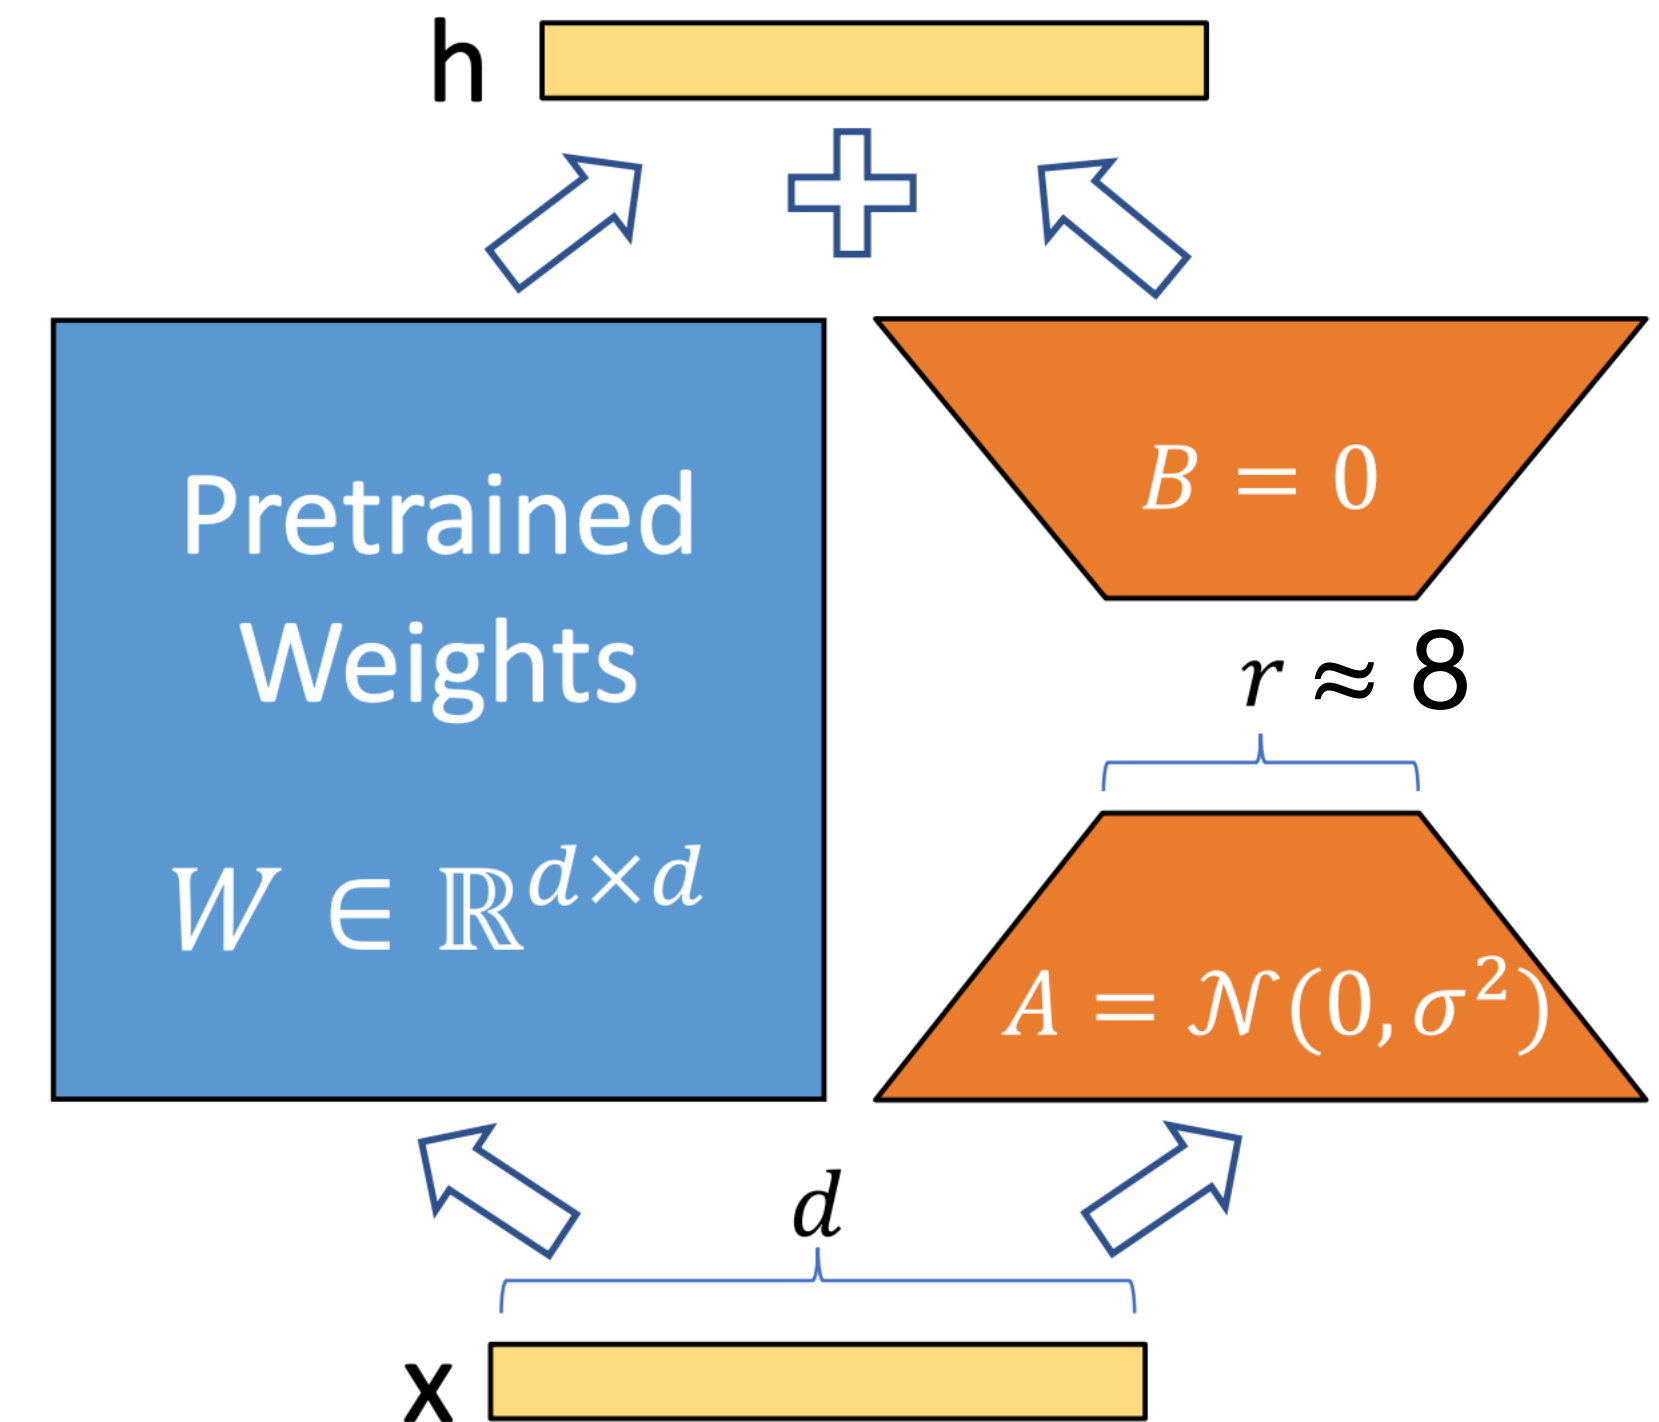
- The basic version of the adapter adds quite a few parameters compared to Prompt Tuning
- Adapters add extra layers that cannot be computed in parallel
- This slows down the model. Especially for small batch sizes

# LoRA

Low-Rank Adaptation

- **Idea:** To adapt the model to the new task, we need to shift the weights towards the anti-gradient direction

$$W' = W + \delta W$$



# LoRA

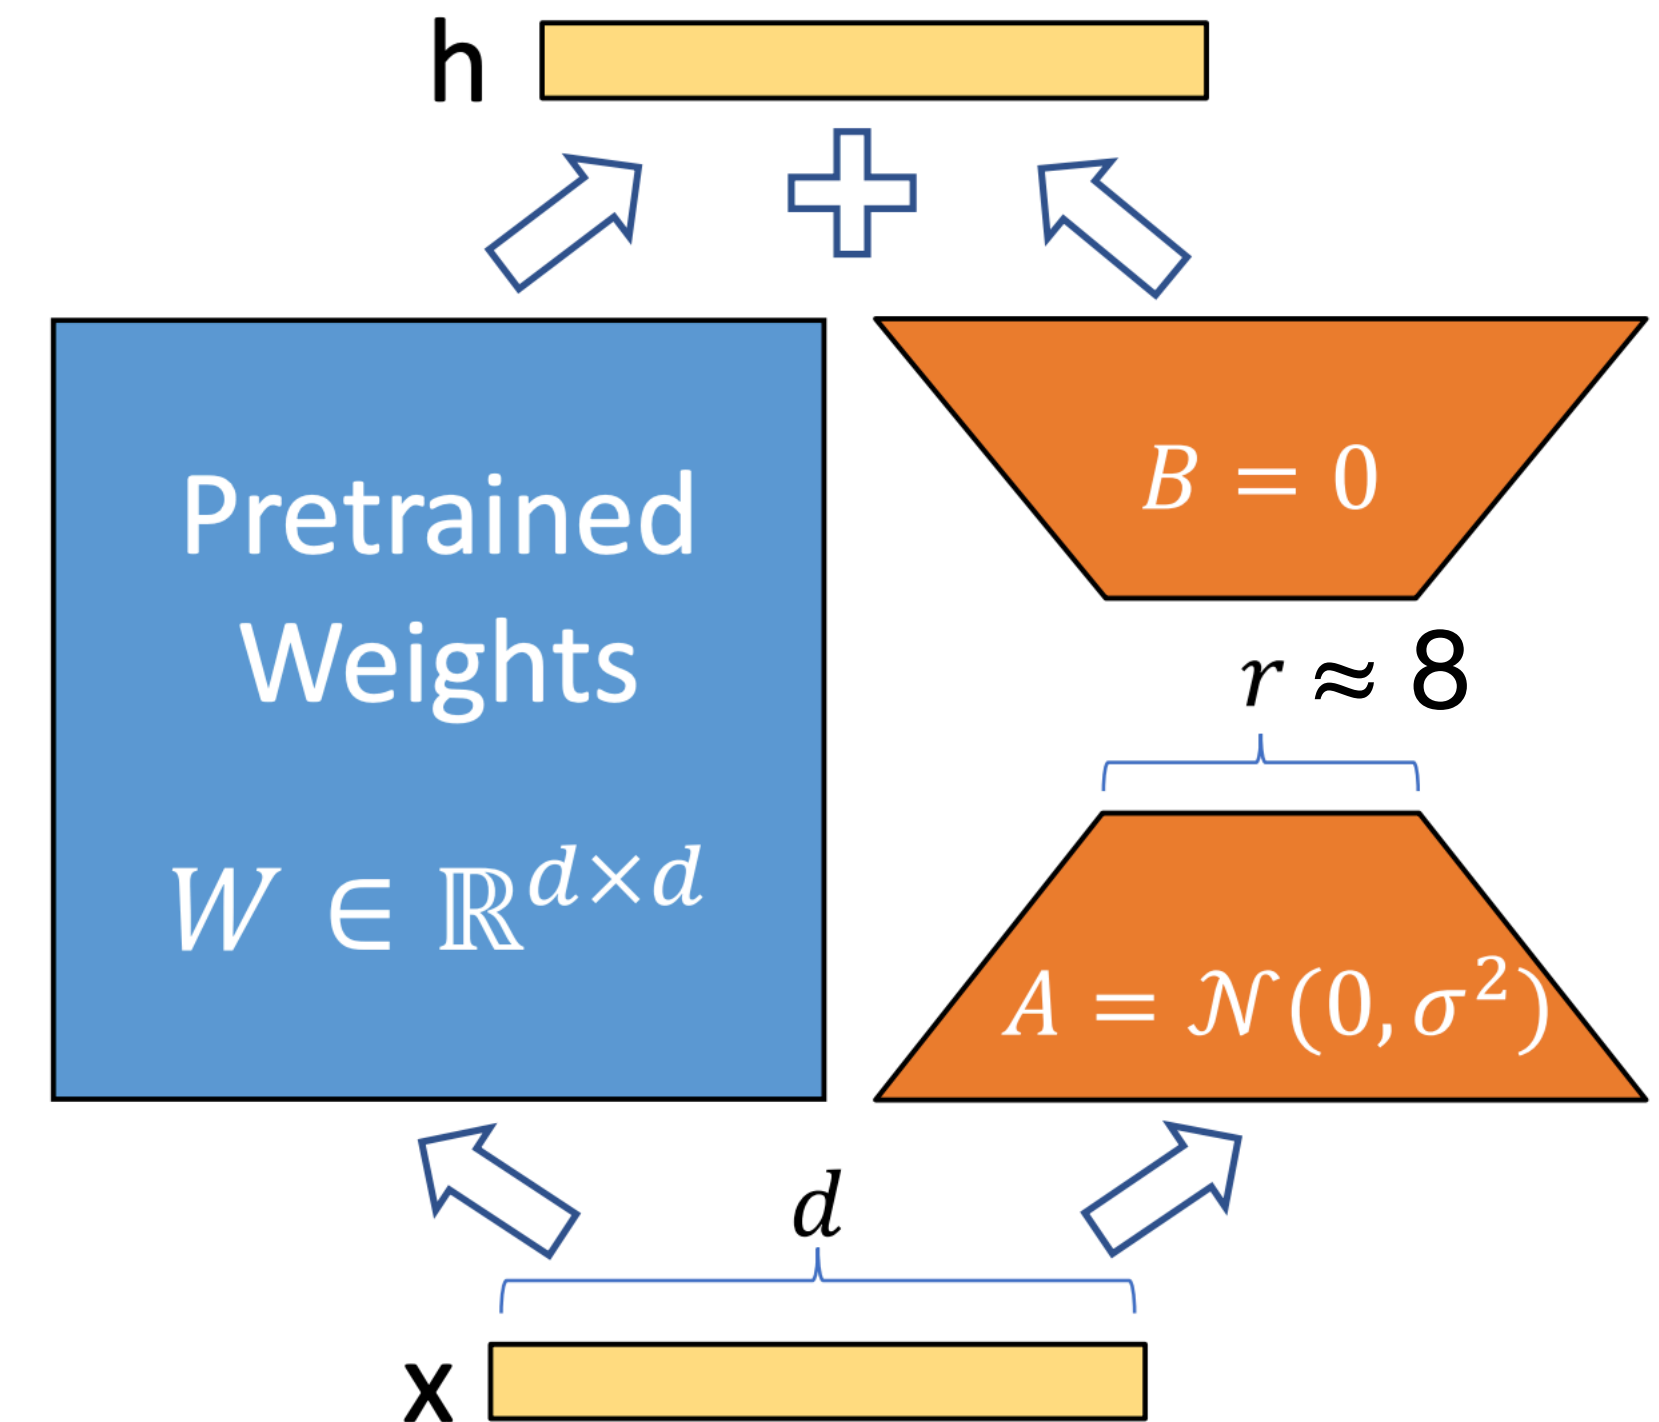
## Low-Rank Adaptation

- **Idea:** To adapt the model to the new task, we need to shift the weights towards the anti-gradient direction

$$W' = W + \delta W$$

- We approximate  $\delta W$  by the product of the training matrices  $AB$

$$W' = W + AB$$



# LoRA

## Low-Rank Adaptation

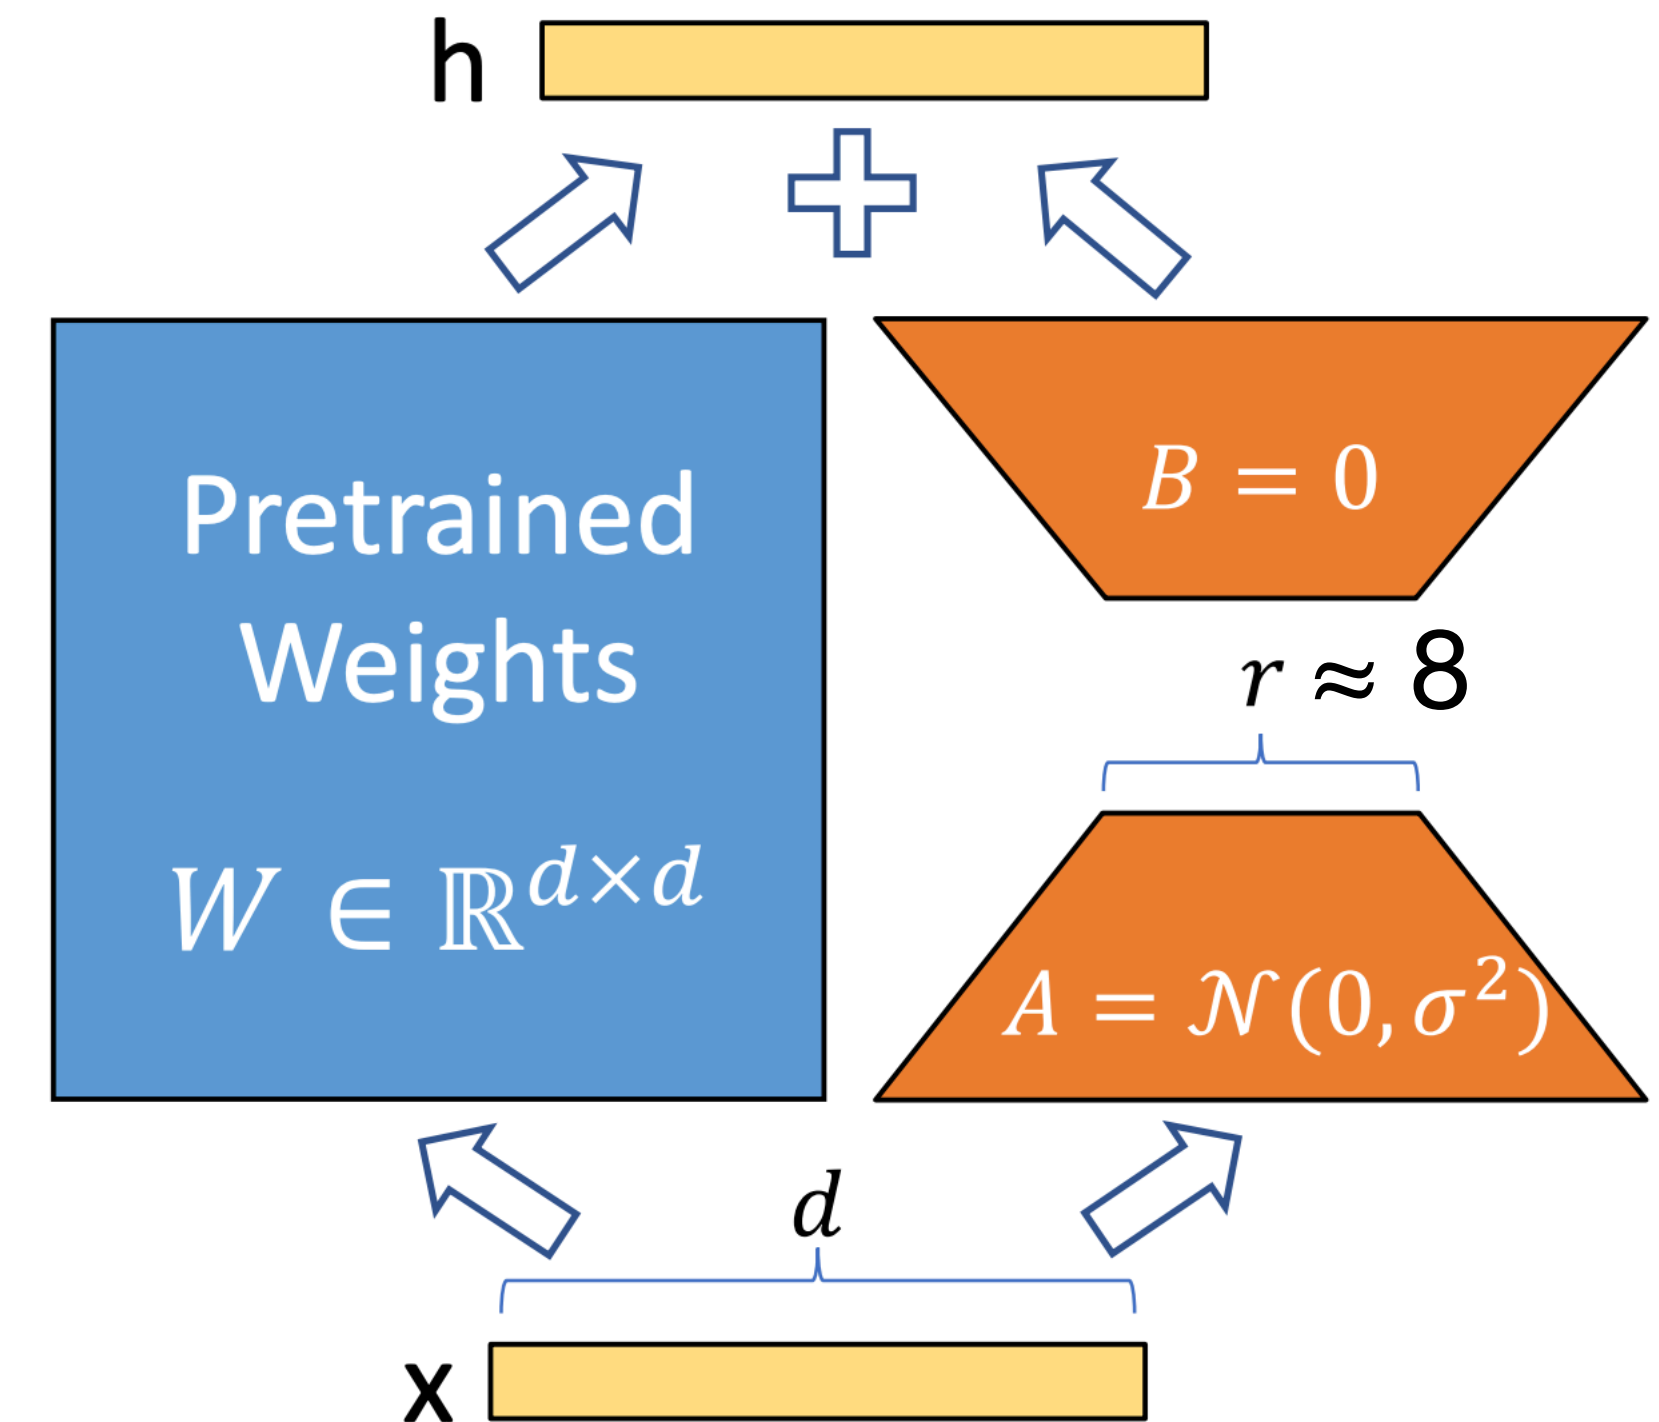
- **Idea:** To adapt the model to the new task, we need to shift the weights towards the anti-gradient direction

$$W' = W + \delta W$$

- We approximate  $\delta W$  by the product of the training matrices  $AB$

$$W' = W + AB$$

- Only the matrices  $W_q$  and  $W_v$  of the attention mechanism are changed in this way
- This way very few parameters are added and the addition can be considered in parallel with the main block
- The most popular PEFT method



# BitFit

## Bias-terms Fine-tuning

**Idea:** biases have very few parameters, we will train only them

### Attention

$$\mathbf{Q}^{m,\ell}(\mathbf{x}) = \mathbf{W}_q^{m,\ell} \mathbf{x} + \mathbf{b}_q^{m,\ell}$$

$$\mathbf{K}^{m,\ell}(\mathbf{x}) = \mathbf{W}_k^{m,\ell} \mathbf{x} + \mathbf{b}_k^{m,\ell}$$

$$\mathbf{V}^{m,\ell}(\mathbf{x}) = \mathbf{W}_v^{m,\ell} \mathbf{x} + \mathbf{b}_v^{m,\ell}$$

$$\mathbf{h}_1^\ell = \text{att}(\mathbf{Q}^{1,\ell}, \mathbf{K}^{1,\ell}, \mathbf{V}^{1,\ell}, \dots, \mathbf{Q}^{m,\ell}, \mathbf{K}^{m,\ell}, \mathbf{V}^{m,\ell})$$

### Feed Forward Network

$$\mathbf{h}_2^\ell = \text{Dropout}(\mathbf{W}_{m_1}^\ell \cdot \mathbf{h}_1^\ell + \mathbf{b}_{m_1}^\ell)$$

$$\mathbf{h}_3^\ell = \mathbf{g}_{LN_1}^\ell \odot \frac{(\mathbf{h}_2^\ell + \mathbf{x}) - \mu}{\sigma} + \mathbf{b}_{LN_1}^\ell$$

$$\mathbf{h}_4^\ell = \text{GELU}(\mathbf{W}_{m_2}^\ell \cdot \mathbf{h}_3^\ell + \mathbf{b}_{m_2}^\ell)$$

$$\mathbf{h}_5^\ell = \text{Dropout}(\mathbf{W}_{m_3}^\ell \cdot \mathbf{h}_4^\ell + \mathbf{b}_{m_3}^\ell)$$

$$\text{out}^\ell = \mathbf{g}_{LN_2}^\ell \odot \frac{(\mathbf{h}_5^\ell + \mathbf{h}_3^\ell) - \mu}{\sigma} + \mathbf{b}_{LN_2}^\ell$$

# Comparison of methods

Suppose we have a transformer with 350M parameters and 24 layers.

Method	Number of parameters
Prompt tuning (length: 20)	15k (0.006%)
Adapters (dim 64)	6M (2%)
LoRA (dim 8)	0.8M (0.22%)
BitFit	0.32M (0.09%)

All methods show comparable quality on some tasks, but LoRA is the most stable and is used more often than the others



# Comparison of methods

Model & Method	# Trainable Parameters	MNLI	SST-2	MRPC	CoLA	QNLI	QQP	RTE	STS-B	Avg.
RoB <sub>base</sub> (FT)*	125.0M	<b>87.6</b>	94.8	90.2	<b>63.6</b>	92.8	<b>91.9</b>	78.7	91.2	86.4
RoB <sub>base</sub> (BitFit)*	0.1M	84.7	93.7	<b>92.7</b>	62.0	91.8	84.0	81.5	90.8	85.2
RoB <sub>base</sub> (Adpt <sup>D</sup> )*	0.3M	87.1 $\pm$ .0	94.2 $\pm$ .1	88.5 $\pm$ 1.1	60.8 $\pm$ .4	93.1 $\pm$ .1	90.2 $\pm$ .0	71.5 $\pm$ 2.7	89.7 $\pm$ .3	84.4
RoB <sub>base</sub> (Adpt <sup>D</sup> )*	0.9M	87.3 $\pm$ .1	94.7 $\pm$ .3	88.4 $\pm$ .1	62.6 $\pm$ .9	93.0 $\pm$ .2	90.6 $\pm$ .0	75.9 $\pm$ 2.2	90.3 $\pm$ .1	85.4
RoB <sub>base</sub> (LoRA)	0.3M	87.5 $\pm$ .3	<b>95.1<math>\pm</math>.2</b>	89.7 $\pm$ .7	63.4 $\pm$ 1.2	<b>93.3<math>\pm</math>.3</b>	90.8 $\pm$ .1	<b>86.6<math>\pm</math>.7</b>	<b>91.5<math>\pm</math>.2</b>	<b>87.2</b>

Model & Method	# Trainable Parameters	E2E NLG Challenge				
		BLEU	NIST	MET	ROUGE-L	CIDEr
GPT-2 M (FT)*	354.92M	68.2	8.62	46.2	71.0	2.47
GPT-2 M (Adapter <sup>L</sup> )*	0.37M	66.3	8.41	45.0	69.8	2.40
GPT-2 M (Adapter <sup>L</sup> )*	11.09M	68.9	8.71	46.1	71.3	2.47
GPT-2 M (Adapter <sup>H</sup> )	11.09M	67.3 $\pm$ .6	8.50 $\pm$ .07	46.0 $\pm$ .2	70.7 $\pm$ .2	2.44 $\pm$ .01
GPT-2 M (FT <sup>Top2</sup> )*	25.19M	68.1	8.59	46.0	70.8	2.41
GPT-2 M (PreLayer)*	0.35M	69.7	8.81	46.1	71.4	2.49
GPT-2 M (LoRA)	0.35M	<b>70.4<math>\pm</math>.1</b>	<b>8.85<math>\pm</math>.02</b>	<b>46.8<math>\pm</math>.2</b>	<b>71.8<math>\pm</math>.1</b>	<b>2.53<math>\pm</math>.02</b>